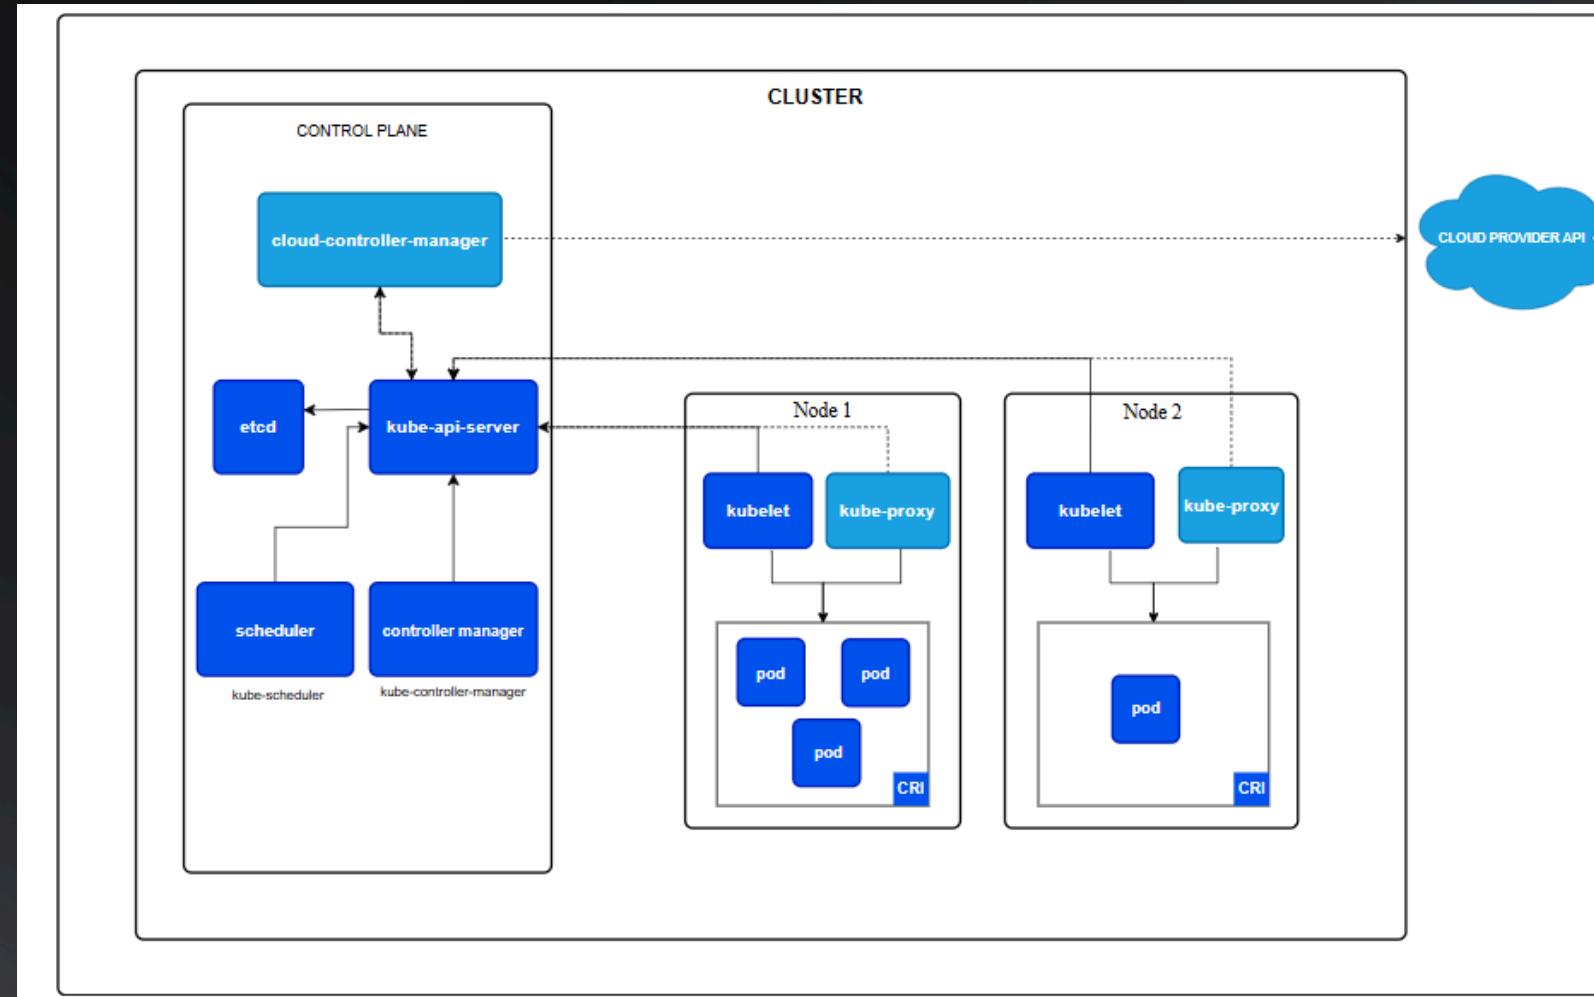


FINAL PROJECT : DATS 6103

Predicting Kubernetes Pod Overload Using Machine Learning

ADITI SHUKLA, ASHLEY GYAPOMAH

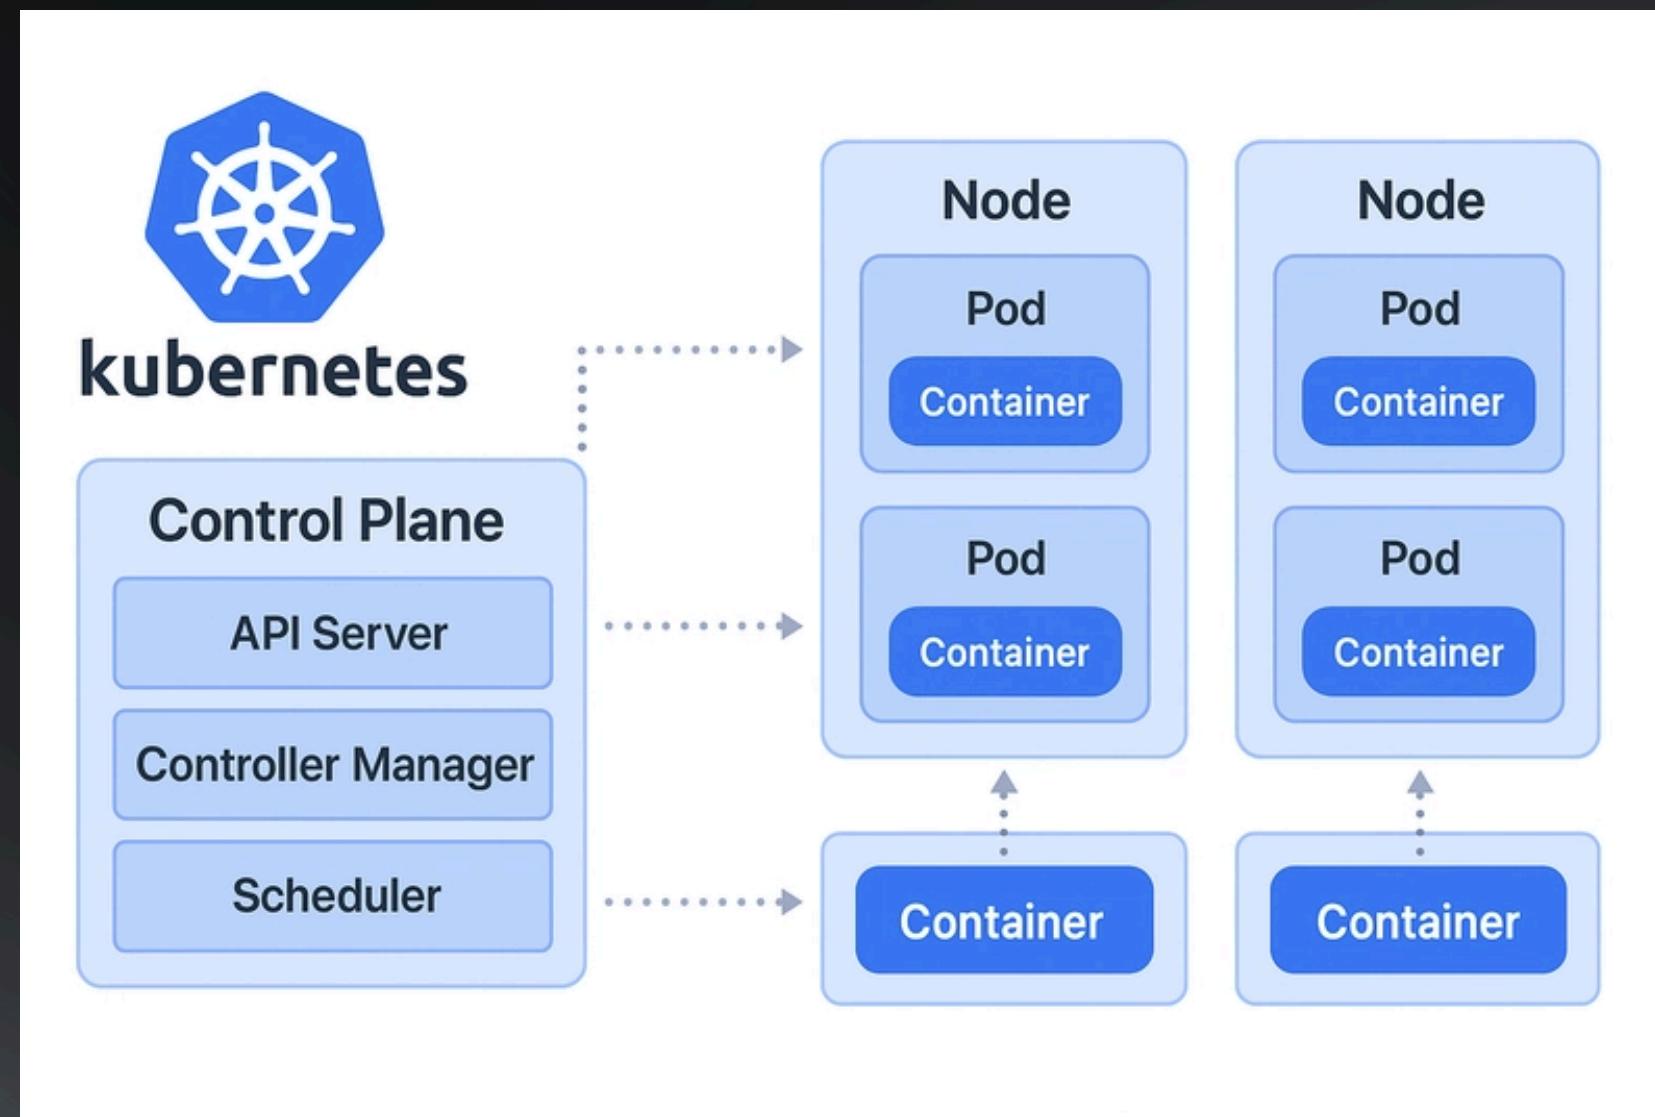
WHAT IS KUBERNATES?



- **KUBERNETES IS AN APPLICATION USED TO ORCHESTRATE AND MANAGE CONTAINERS ACROSS MANY MACHINES.**
- **EXAMPLE NETFLIX SCALING AUTOMATICALLY WHEN USERS ARE HIGHER**

KUBERNATES ARCHITECTURE

CORE CONCEPT: CONTAINER ORCHESTRATION



WHAT IS A CONTAINER?

- **CONTAINS AND APPLICATIONS CODE, RUNTIME, SYSTEM LIBRARIES AND SETTINGS**

WHAT IS ORCHESTRATION?

- **THE PROCESS OF AUTOMATICALLY ORGANIZING AND MANAGING THOUSANDS OF CONTAINERS**

PROBLEM STATEMENT AND SMART QUESTION

**CAN WE BUILD A MACHINE LEARNING MODEL THAT PREDICTS POD OVERLOAD
(NEED_NEW_POD = 1) USING ONLY RESOURCE ALLOCATION METRICS?**

CURRENT SYSTEM AUTOMATICALLY CREATES A NEW POD ONLY AFTER THE OVERLOAD HAPPENS.
PROBLEM: THIS REACTION IS SLIGHTLY DELAYED → USERS EXPERIENCE LAG/BUFFERING.
GOAL: PREDICT OVERLOAD BEFORE IT HAPPENS SO SCALING CAN OCCUR PROACTIVELY.

SPECIFIC → CPU & MEMORY METRICS
MEASURABLE → F1, ACCURACY, ROC-AUC
ACHIEVABLE → WE ACHIEVED > 95% F1
RELEVANT → DIRECTLY RELATES TO KUBERNETES AUTOSCALING
TIME-BOUND → COMPLETED WITHIN THE PROJECT TIMELINE

DATASET OVERVIEW

DATASET: KUBERNETES RESOURCE ALLOCATION DATASET FROM KAGGLE

```
0    pod_name
1    namespace
2    cpu_request
3    cpu_limit
4    memory_request
5    memory_limit
6    cpu_usage
7    memory_usage
8    node_name
9    pod_status
10   restart_count
11   uptime_seconds
12   deployment_strategy
13   scaling_policy
14   network_bandwidth_usage
```

RESOURCE METRICS

- CPU_REQUEST – CPU REQUESTED BY THE POD
- CPU_LIMIT – MAX CPU ALLOWED
- CPU_USAGE – ACTUAL CPU USED
- MEMORY_REQUEST – MEMORY REQUESTED
- MEMORY_LIMIT – MAX MEMORY ALLOWED
- MEMORY_USAGE – ACTUAL MEMORY USED

POD METADATA

- POD_NAME
- NAMESPACE
- NODE_NAME
- POD_STATUS
- DEPLOYMENT_STRATEGY
- SCALING_POLICY
- OTHER SYSTEM FEATURES

RESTART_COUNT

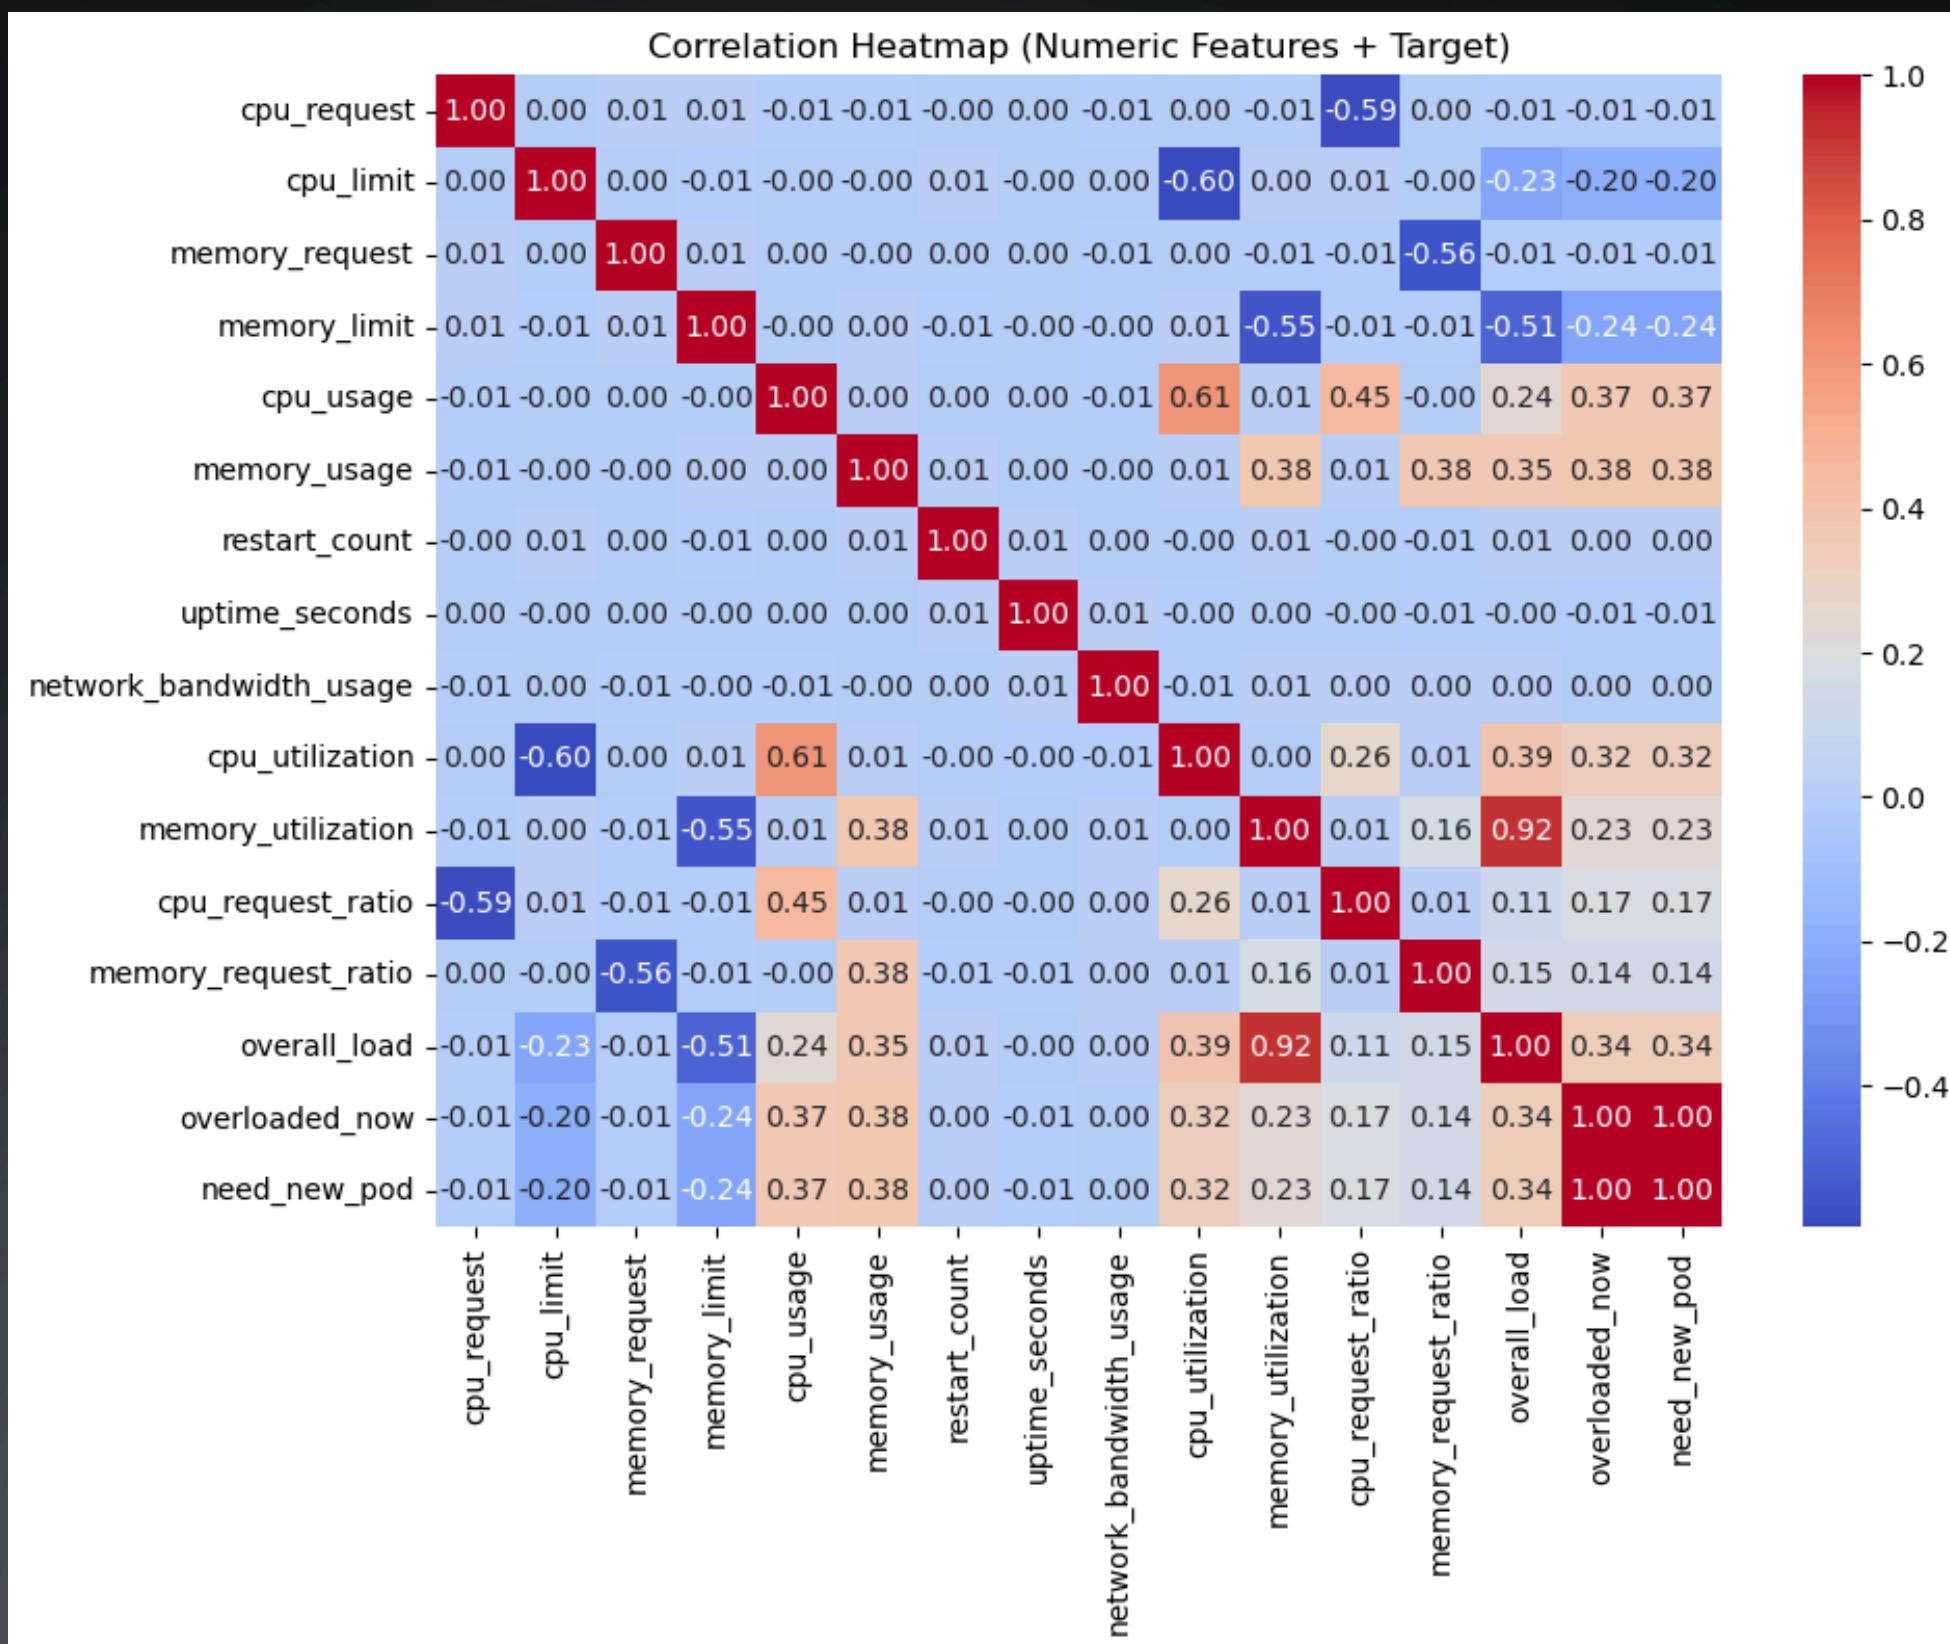
- UPTIME_SECONDS
- NETWORK_BANDWIDTH_USAGE

DATA PREPROCESSING & CLEANING

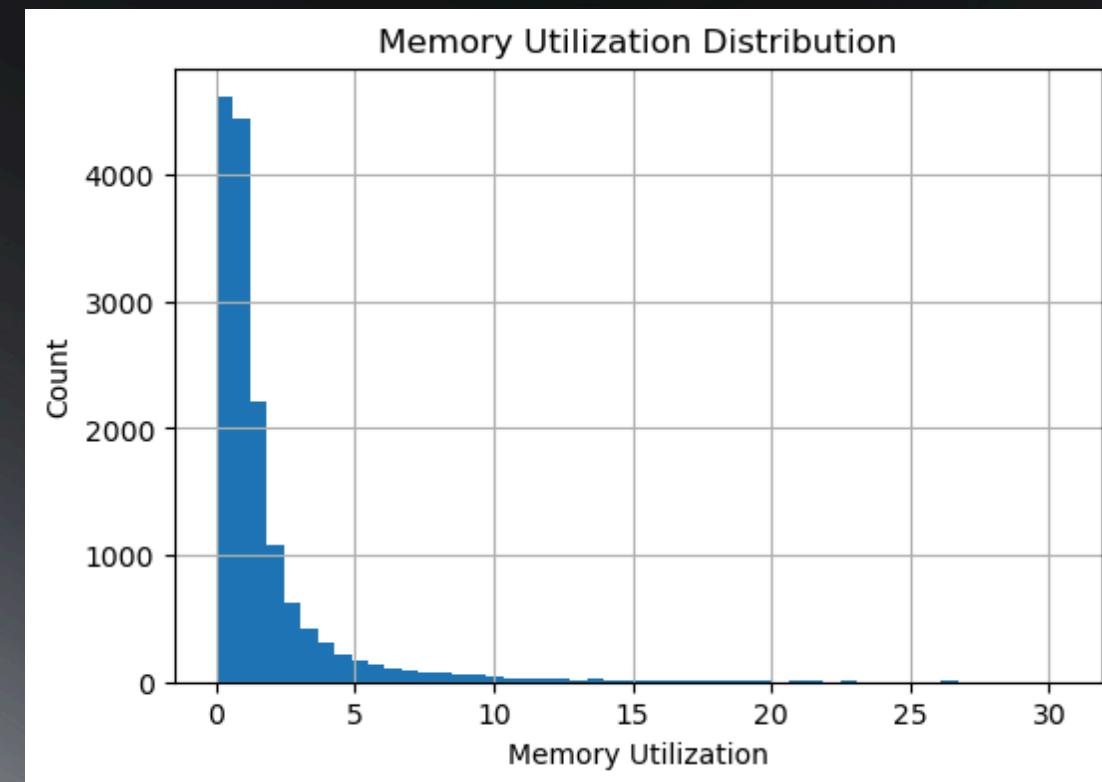
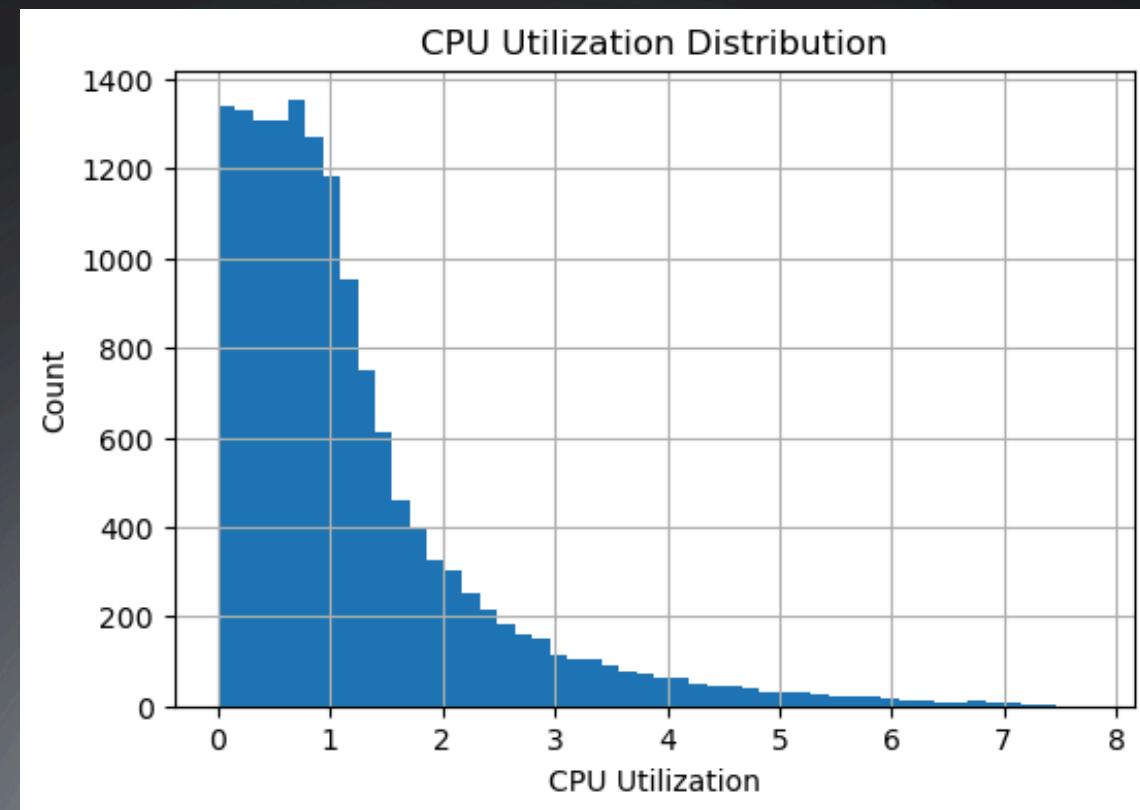
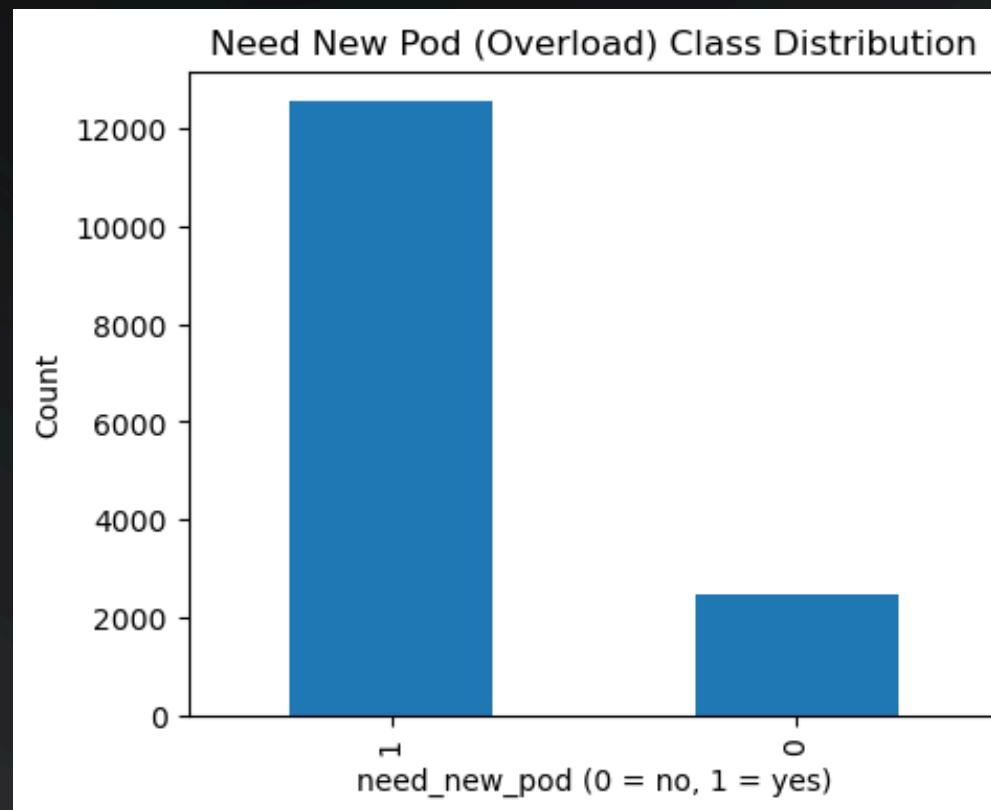
```
100%|██████████| 15000/15000 [00:00<00:00, 149.80it/s]
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   pod_name         15000 non-null   object  
 1   namespace        15000 non-null   object  
 2   cpu_request      15000 non-null   float64 
 3   cpu_limit        15000 non-null   float64 
 4   memory_request   15000 non-null   float64 
 5   memory_limit     15000 non-null   float64 
 6   cpu_usage        15000 non-null   float64 
 7   memory_usage     15000 non-null   float64 
 8   node_name        15000 non-null   object  
 9   pod_status       15000 non-null   object  
 10  restart_count    15000 non-null   int64  
 11  uptime_seconds  15000 non-null   int64  
 12  deployment_strategy 15000 non-null   object  
 13  scaling_policy   15000 non-null   object  
 14  network_bandwidth_usage 15000 non-null   float64 
dtypes: float64(7), int64(2), object(6)
...
deployment_strategy      0
scaling_policy           0
network_bandwidth_usage 0
dtype: int64
```

- ALL COLUMNS HAVE 15,000 NON-NULL ENTRIES
- 0 MISSING VALUES
- CHECKED FOR NEGATIVE VALUES IN COLUMNS
BECAUSE IT WOULD BE INVALID
- NORMALIZED COLUMN NAMES TO LOWER CASE
- DATASET IS CLEAN AND COMPLETE

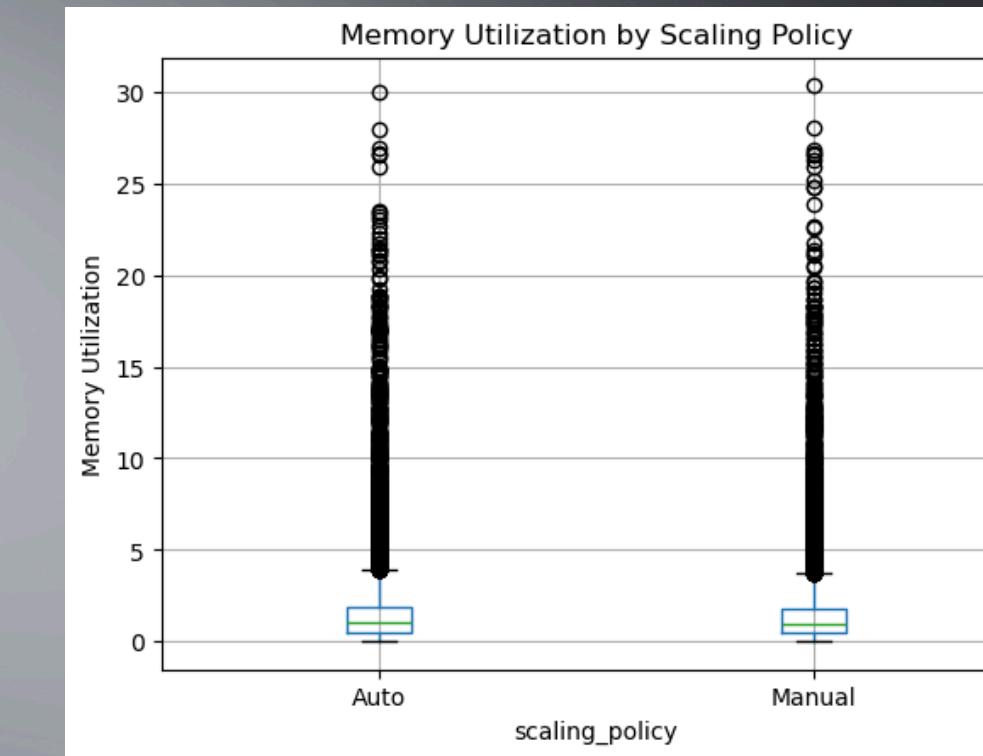
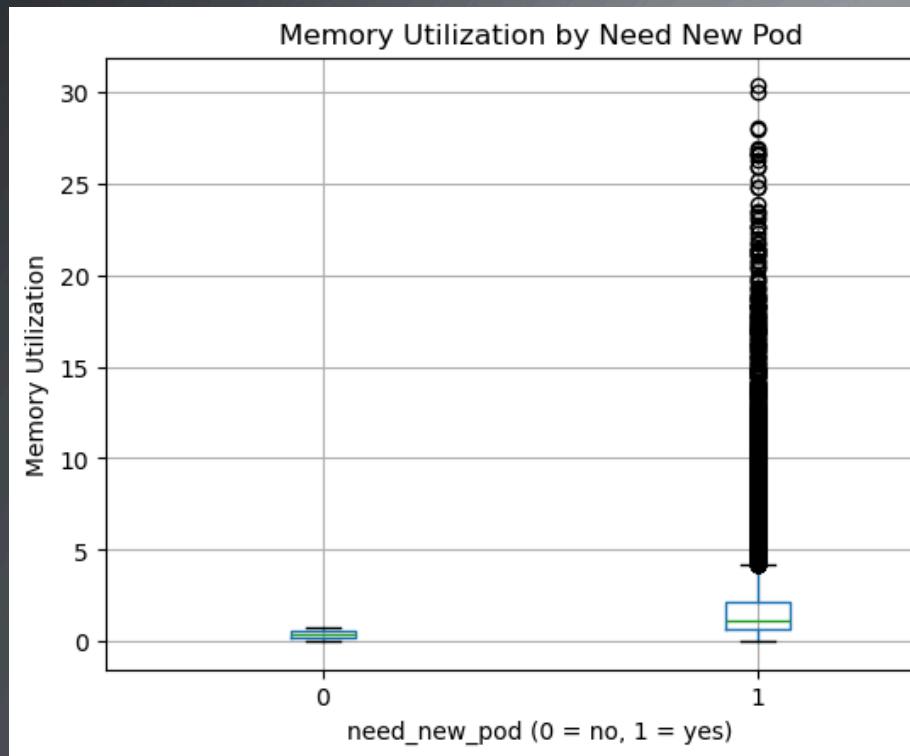
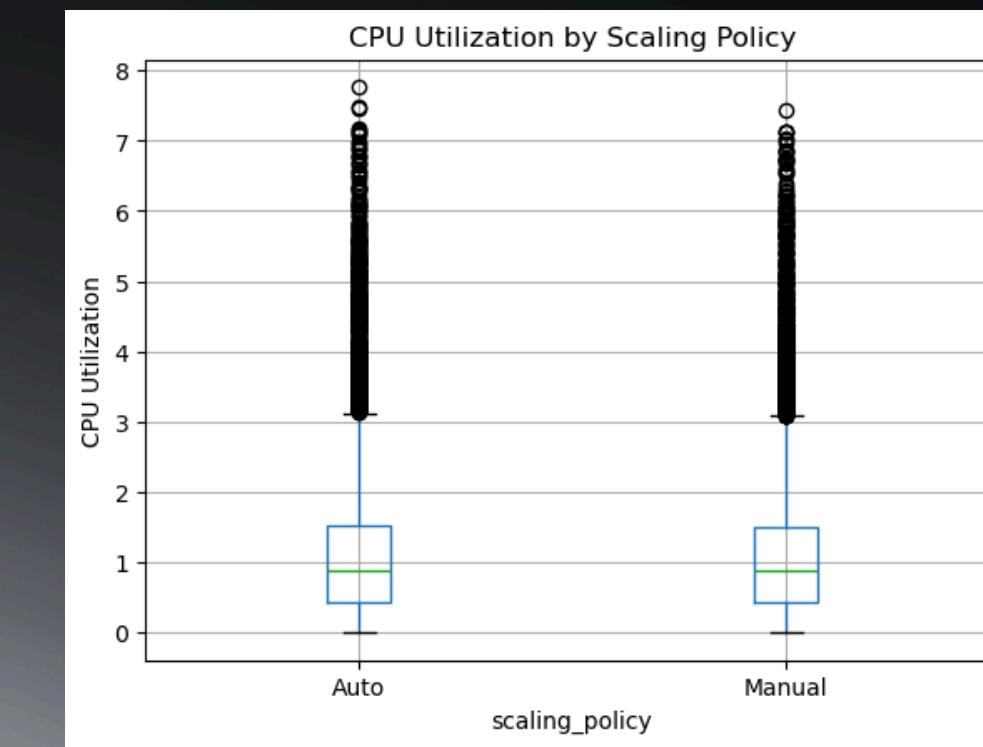
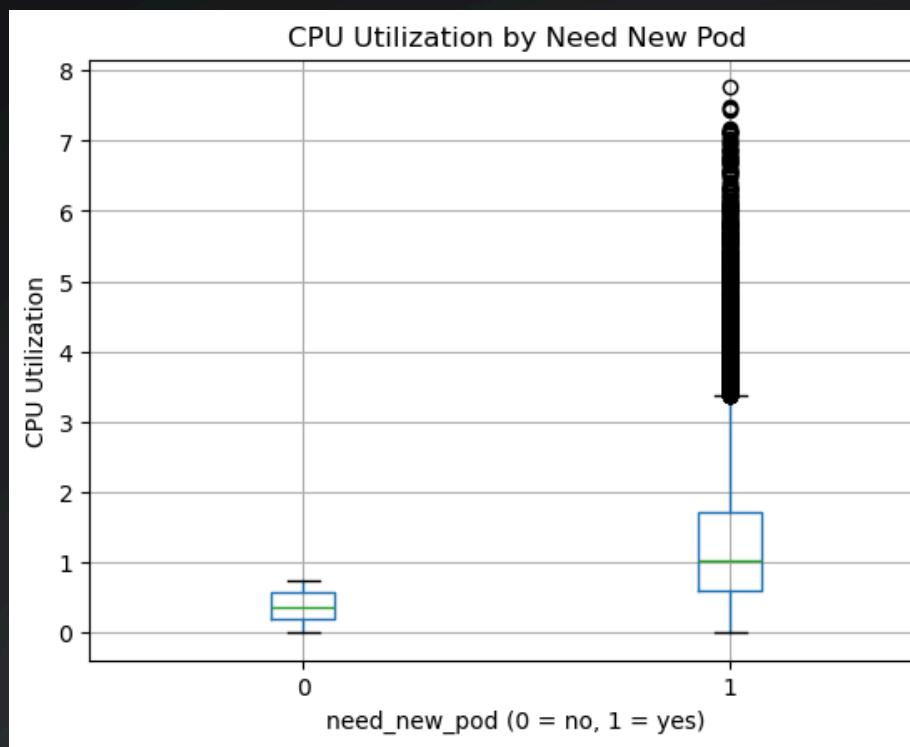
EXPLORATORY DATA ANALYSIS



EXPLORATORY DATA ANALYSIS



EXPLORATORY DATA ANALYSIS



STATISTICAL TESTING

ANOVA F-statistic: 0.7010766503854831

ANOVA P-value: 0.5512752358978207

No significant difference between groups.

T-statistic: 65.85529426785459

P-value: 0.0

The means are significantly different.

Chi-square: 3.5775707777555152

P-value: 0.05856473401330775

Variables are independent.

FEATURE ENGINEERING

- CPU & MEMORY UTILIZATION ($\text{USAGE} \div \text{LIMIT}$) : ROW 4 → $\text{CPU_UTILIZATION} = 1.67 \rightarrow \text{CPU USAGE EXCEEDS THE ALLOWED LIMIT} \rightarrow \text{RISK OF OVERLOAD}$
- USAGE VS REQUEST RATIO ($\text{USAGE} \div \text{REQUEST}$) : ROW 3 → $\text{CPU_REQUEST_RATIO} = 12.04 \rightarrow \text{USED 12x MORE CPU THAN REQUESTED} \rightarrow \text{VERY UNSTABLE POD}$
- COMBINED LOAD METRIC (OVERALL LOAD): ROW 1 → $\text{OVERALL_LOAD} = 1.37 \rightarrow \text{HEAVY LOAD, LIKELY TO REQUIRE A NEW POD SOON.}$

	<code>cpu_request</code>	<code>cpu_limit</code>	<code>cpu_usage</code>	<code>memory_request</code>	<code>memory_limit</code>	<code>memory_usage</code>
0	1.569542	3.679152	3.345496	3174.582783	5134.413852	2135.310365
1	0.343119	3.722716	2.758188	3551.459173	3698.349366	7442.200271
2	0.249271	1.318147	1.319703	1578.313253	7418.271122	5142.897754
3	0.311497	2.852595	3.752312	1392.962372	3628.480705	2952.449331
4	1.532775	0.521618	0.874224	2660.192655	5091.497752	3382.299355

<code>cpu_utilization</code>	<code>memory_utilization</code>	<code>cpu_request_ratio</code>	<code>memory_request_ratio</code>	<code>overall_load</code>
0.909311	0.415882	2.131510	0.672627	0.662597
0.740907	2.012303	8.038545	2.095533	1.376605
1.001180	0.693274	5.294233	3.258477	0.847227
1.315403	0.813687	12.046042	2.119547	1.064545
1.675984	0.664303	0.570353	1.271449	1.170143

TARGET VARIABLE

Target Variable: scaling_event

Binary Target (0/1):

- 1 → Scaling needed (pod was overloaded / near limits)
- 0 → No scaling needed (pod running normally)

```
# Thresholds for overload
cpu_thresh = 0.75
mem_thresh = 0.75
```

MODEL TRAINING AND EVALUATION

```
Final modeling dataset shape: (15000, 19) (15000,)
```

```
Target distribution:
```

```
need_new_pod
```

```
1    12552
```

```
0     2448
```

```
Name: count, dtype: int64
```

```
Train/Test Split Complete
```

```
Training samples: 12000
```

```
Testing samples : 3000
```

Feature & Target Selection

- Selected key numeric features (CPU/memory usage, limits, ratios, load)
- Selected categorical features (namespace, status, scaling policy)

Final Modeling Dataset

- No missing values
- X = selected features
- y = need_new_pod

Train/Test Split

- 80% train, 20% test
- Stratified to preserve target class ratio
- 12,000 training samples
- 3,000 testing samples

- Scaled all the numeric values
- used ‘onehotencoder’ to convert categorical values in to numeric
- combined it using ‘columntransformer’

MODEL TRAINING AND EVALUATION

```
... ===== Training LogisticRegression =====
Accuracy : 0.9343
Precision: 0.9591
Recall   : 0.9625
F1-score : 0.9608
ROC-AUC  : 0.9796

Classification report:
precision    recall  f1-score   support
          0       0.80      0.79      0.80      490
          1       0.96      0.96      0.96     2510
accuracy                           0.93      3000
macro avg       0.88      0.88      0.88      3000
weighted avg    0.93      0.93      0.93      3000
```

```
===== Training DecisionTree =====
Accuracy : 0.9983
Precision: 0.9984
Recall   : 0.9996
F1-score : 0.999
ROC-AUC  : 0.9969

Classification report:
precision    recall  f1-score   support
          0       1.00      0.99      0.99      490
          1       1.00      1.00      1.00     2510
accuracy                           1.00      3000
macro avg       1.00      1.00      1.00      3000
weighted avg    1.00      1.00      1.00      3000
```

MODEL TRAINING AND EVALUATION

```
===== Training RandomForest =====
```

```
Accuracy : 0.9987  
Precision: 0.9984  
Recall   : 1.0  
F1-score : 0.9992  
ROC-AUC  : 1.0
```

```
Classification report:
```

	precision	recall	f1-score	support
0	1.00	0.99	1.00	490
1	1.00	1.00	1.00	2510
accuracy			1.00	3000
macro avg	1.00	1.00	1.00	3000
weighted avg	1.00	1.00	1.00	3000

```
===== Training GradientBoosting =====
```

```
Accuracy : 0.9997  
Precision: 0.9996  
Recall   : 1.0  
F1-score : 0.9998  
ROC-AUC  : 0.9999
```

```
Classification report:
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	490
1	1.00	1.00	1.00	2510
accuracy			1.00	3000
macro avg	1.00	1.00	1.00	3000
weighted avg	1.00	1.00	1.00	3000

MODEL TRAINING AND EVALUATION

```
===== Training XGBoost =====
Accuracy : 0.9993
Precision: 0.9996
Recall   : 0.9996
F1-score : 0.9996
ROC-AUC  : 1.0

classification report:
      precision    recall  f1-score   support
          0         1.00     1.00     1.00      490
          1         1.00     1.00     1.00    2510
  accuracy                           1.00      3000
  macro avg       1.00     1.00     1.00      3000
weighted avg       1.00     1.00     1.00      3000
```

MODEL TRAINING AND EVALUATION

